

Data Intensive Computing - Review Questions 1

Deadline: September 12, 2021

1. **1 point.** Explain how a file region can be left in an **undefined** state in GFS?

2. **1 point.** Briefly explain the read operation in GFS? What's the role of the primary node?

3. **1 point.** Using one example show that in the CAP theorem, if we have Consistency and Partition Tolerance, we cannot provide Availability at the same time.

4. **1 point.** Explain how the consistent hashing works.

5. **1 point.** Explain the finding process of tablets in BigTable.